

```
In [1]: from keras.models import Sequential
        from keras.layers import Dense

        import numpy as np

        import matplotlib.pyplot as plt

        %matplotlib inline

        from tensorflow.examples.tutorials.mnist import input_data

        mnist = input_data.read_data_sets('./', one_hot=True)
```

Using TensorFlow backend.


```
Successfully downloaded train-images-idx3-ubyte.gz 9912422 bytes.
Extracting ./train-images-idx3-ubyte.gz
Successfully downloaded train-labels-idx1-ubyte.gz 28881 bytes.
Extracting ./train-labels-idx1-ubyte.gz
Successfully downloaded t10k-images-idx3-ubyte.gz 1648877 bytes.
Extracting ./t10k-images-idx3-ubyte.gz
Successfully downloaded t10k-labels-idx1-ubyte.gz 4542 bytes.
Extracting ./t10k-labels-idx1-ubyte.gz
```

```
In [10]: x, y = mnist.train.next_batch(20)
        x.shape
```

Out[10]: (20, 784)

```
In [3]: model = Sequential() #initiates a layer
        model.add(Dense(100, activation='relu', input_dim=784)) # adds a layer
        model.add(Dense(10, activation='softmax')) #one layer of softmax
            #required
        model.compile(loss='categorical_crossentropy',
                      optimizer='adam',
                      metrics=['accuracy'])
```

```
In [19]: plt.imshow(model.get_weights()[0][:,40].reshape(28,28))  
plt.show()
```



```
In [4]: def get_batch(dataset, batch_size = 256):  
        while (1):  
            yield dataset.next_batch(batch_size)  
            #continues where left off if called again(new  
            #test set trials are different)
```

```
In [5]: batch_size = 256  
test_gen = get_batch(mnist.test, batch_size)  
steps_per_epoch = mnist.test.num_examples// batch_size  
model.evaluate_generator(test_gen, steps_per_epoch)  
#untrained model
```

```
Out[5]: [2.4107079261388535, 0.12770432692307693]
```

```
In [11]: batch_size = 256
data_gen = get_batch(mnist.train, batch_size)
steps_per_epoch = mnist.train.num_examples//batch_size

model.fit_generator(data_gen, steps_per_epoch, epochs=10)
#trains model to data

Epoch 1/10
214/214 [=====] - 1s - loss: 0.0541 - acc: 0.9847
Epoch 2/10
214/214 [=====] - 1s - loss: 0.0490 - acc: 0.9864
Epoch 3/10
214/214 [=====] - 1s - loss: 0.0430 - acc: 0.9882
Epoch 4/10
214/214 [=====] - 1s - loss: 0.0391 - acc: 0.9894
Epoch 5/10
214/214 [=====] - 1s - loss: 0.0351 - acc: 0.9904
Epoch 6/10
214/214 [=====] - 1s - loss: 0.0317 - acc: 0.9919
Epoch 7/10
214/214 [=====] - 1s - loss: 0.0285 - acc: 0.9928
Epoch 8/10
214/214 [=====] - 1s - loss: 0.0255 - acc: 0.9943
Epoch 9/10
214/214 [=====] - 1s - loss: 0.0229 - acc: 0.9949
Epoch 10/10
214/214 [=====] - 1s - loss: 0.0213 - acc: 0.9951
```

```
Out[11]: <keras.callbacks.History at 0x7ffa9411b8d0>
```

```
In [7]: model.evaluate_generator(test_gen, steps_per_epoch)
```

```
Out[7]: [0.085014622623675332, 0.97282053154205606]
```

It is really important to be able to reload the model after you've been training it for hours on end (usually).

```
In [8]: from keras.models import load_model  
  
        model.save('my_model.h5')  
  
        model2 = load_model('my_model.h5')
```

```
In [9]: model2.evaluate_generator(test_gen, steps_per_epoch)
```

```
Out[9]: [0.084644120517317387, 0.97276577102803741]
```

```
In [ ]:
```